



Fast Detection of Near Duplicates

Sambuddha Roy, LinkedIn

July 2, 2016



Motivation.



and some more...



and more...



The duplicate detection problem.

- ▶ Why are duplicates problematic?



The duplicate detection problem.

- ▶ Why are duplicates problematic?
- ▶ Imagine if...



The duplicate detection problem.

- ▶ Why are duplicates problematic?
- ▶ Imagine if...
- ▶ Youtube had multiple copies of each video!
 - ▶ Someone copies the cat video that *you* uploaded!



The duplicate detection problem.

- ▶ Why are duplicates problematic?
- ▶ Imagine if...
- ▶ Youtube had multiple copies of each video!
 - ▶ Someone copies the cat video that *you* uploaded!
- ▶ Duplicate content confuses search engines.
 - ▶ What happens to the pageRank of the page?



How do you catch duplicates?

- ▶ Some form of hashing...
- ▶ Different items go to different hashes (collision resistant).
- ▶ Group by hashes - each group corresponds to a distinct element.



How about near-duplicates?

► Now imagine if...



How about near-duplicates?

- ▶ Now imagine if...
- ▶ Someone made a profile on Facebook, very similar to yours.
- ▶ copied profile pictures,



How about near-duplicates?

- ▶ Now imagine if...
- ▶ Someone made a profile on Facebook, very similar to yours.
- ▶ copied profile pictures, connected to the same friends,



How about near-duplicates?

- ▶ Now imagine if...
- ▶ Someone made a profile on Facebook, very similar to yours.
- ▶ copied profile pictures, connected to the same friends, indicated same interests



How about near-duplicates?

- ▶ Now imagine if...
- ▶ Someone made a profile on Facebook, very similar to yours.
- ▶ copied profile pictures, connected to the same friends, indicated same interests etc.



How about near-duplicates?

- ▶ Now imagine if...
- ▶ Someone made a profile on Facebook, very similar to yours.
- ▶ copied profile pictures, connected to the same friends, indicated same interests etc.
- ▶ Identity theft!



How about near-duplicates?

- ▶ Now imagine if...
- ▶ Someone made a profile on Facebook, very similar to yours.
- ▶ copied profile pictures, connected to the same friends, indicated same interests etc.
- ▶ Identity theft!
- ▶ Other use cases: plagiarism, etc.



Essential Problem: Nearest Neighbor Search

- ▶ There is a database of items \mathcal{D} .
- ▶ ...and a query q arrives.



Essential Problem: Nearest Neighbor Search

- ▶ There is a database of items \mathcal{D} .
- ▶ ...and a query q arrives.
- ▶ Given the query q , find the nearest neighbors of q in the database \mathcal{D} .



Essential Problem: Nearest Neighbor Search

- ▶ There is a database of items \mathcal{D} .
- ▶ ...and a query q arrives.
- ▶ Given the query q , find the nearest neighbors of q in the database \mathcal{D} .
- ▶ Often, we just want the k nearest neighbors (k-NN problem).



Nearness?

- ▶ Nearness only makes sense in the presence of a *distance measure*.



Nearness?

- ▶ Nearness only makes sense in the presence of a *distance measure*.
- ▶ So, similarity measures between objects/items...
- ▶ i.e. *featurize* items



Nearness?

- ▶ Nearness only makes sense in the presence of a *distance measure*.
- ▶ So, similarity measures between objects/items...
- ▶ i.e. *featurize* items as vectors $\in \mathbb{R}^n$ or in $\{0, 1\}^n$.



Nearness?

- ▶ Nearness only makes sense in the presence of a *distance measure*.
- ▶ So, similarity measures between objects/items...
- ▶ i.e. *featurize* items as vectors $\in \mathbb{R}^n$ or in $\{0, 1\}^n$.
- ▶ And consider some distance/similarity measure between these vectors.



Typical Distance/Similarity Measures

► Jaccard similarity: $J(A, B) = \frac{|A \cap B|}{|A \cup B|}$



Typical Distance/Similarity Measures

- ▶ Jaccard similarity: $J(A, B) = \frac{|A \cap B|}{|A \cup B|}$
- ▶ Cosine Similarity etc.: $\text{sim}(a, b) = \frac{a \cdot b}{|a||b|}$



Typical Distance/Similarity Measures

- ▶ Jaccard similarity: $J(A, B) = \frac{|A \cap B|}{|A \cup B|}$
- ▶ Cosine Similarity etc.: $\text{sim}(a, b) = \frac{a \cdot b}{|a||b|}$
- ▶ Hamming Distance
- ▶ ℓ_p for $p \in (0, 2]$.
- ▶ Tanimoto distance, Mahalanobis distance, etc.



Typical Distance/Similarity Measures

- ▶ Jaccard similarity: $J(A, B) = \frac{|A \cap B|}{|A \cup B|}$
- ▶ Cosine Similarity etc.: $\text{sim}(a, b) = \frac{a \cdot b}{|a||b|}$
- ▶ Hamming Distance
- ▶ ℓ_p for $p \in (0, 2]$.
- ▶ Tanimoto distance, Mahalanobis distance, etc.
- ▶ It's typical to relate the distance and similarity measures as $\text{dist}(a, b) = 1 - \text{sim}(a, b)$.



Back to what we want...

- ▶ Given a *query* object q , we would want to retrieve all the *database* items that are “similar” to q .



Back to what we want...

- ▶ Given a *query* object q , we would want to retrieve all the *database* items that are “similar” to q .
- ▶ A search by pairwise comparisons between q and all the items in the database may become too costly.



Another use-case...

- ▶ If we want to “cluster” the items in the database according to their similarities/distances.



Another use-case...

- ▶ If we want to “cluster” the items in the database according to their similarities/distances.
- ▶ (Aside: what does this even mean? It may be for a threshold τ and points a, b, c that $\text{dist}(a, b) \leq \kappa$ and $\text{dist}(b, c) \leq \kappa$ but $\text{dist}(a, c) > \kappa$ i.e. similarity is not “transitive”).



Similarity is not transitive

- ▶ A more meaningful formulation can be: we want **all** pairs of items a, b such that $\text{sim}(a, b) \geq \tau$ for some $\tau \in [0, 1]$ (sufficiently similar items).



Similarity is not transitive

- ▶ A more meaningful formulation can be: we want **all** pairs of items a, b such that $\text{sim}(a, b) \geq \tau$ for some $\tau \in [0, 1]$ (sufficiently similar items).
- ▶ If we make $\tau = 0$ we are asking for all of the $\binom{n}{2}$ pairs of items (for n items in the database).



Similarity is not transitive

- ▶ A more meaningful formulation can be: we want **all** pairs of items a, b such that $\text{sim}(a, b) \geq \tau$ for some $\tau \in [0, 1]$ (sufficiently similar items).
- ▶ If we make $\tau = 0$ we are asking for all of the $\binom{n}{2}$ pairs of items (for n items in the database).
- ▶ Imagine if the number of items n were



Similarity is not transitive

- ▶ A more meaningful formulation can be: we want **all** pairs of items a, b such that $\text{sim}(a, b) \geq \tau$ for some $\tau \in [0, 1]$ (sufficiently similar items).
- ▶ If we make $\tau = 0$ we are asking for all of the $\binom{n}{2}$ pairs of items (for n items in the database).
- ▶ Imagine if the number of items n were 1 million...



Similarity is not transitive

- ▶ A more meaningful formulation can be: we want **all** pairs of items a, b such that $\text{sim}(a, b) \geq \tau$ for some $\tau \in [0, 1]$ (sufficiently similar items).
- ▶ If we make $\tau = 0$ we are asking for all of the $\binom{n}{2}$ pairs of items (for n items in the database).
- ▶ Imagine if the number of items n were 1 million... 1 billion...!



Alternate Formulation/Relaxations

- ▶ We really do not want **all** pairs for the similarity threshold, τ really small; in fact, typical use-cases will consider $\tau > 0.8$ or so (sufficiently similar).
- ▶ Let's also relax the **all** in the above too; replace that by **most**.



Refining what we want

- ▶ Find near-duplicates of query items.
- ▶ Some mistakes will be allowed (both false positives, false negatives).
- ▶ Time! Querying should be **fast**. The clustering variant should take $O(n)$ time instead of $O(n^2)$.



Hashing to the rescue

- ▶ Can we hash items so that “nearby” items are in **same** hash-buckets?
- ▶ Note that this is counter to the usual notion of hashing, where collisions are taboo.
- ▶ Here, we would like collisions - but only between *nearby* items.



Enter LSH

- ▶ Locality Sensitive Hashing - introduced by Indyk & Motwani, in 1998.
 - ▶ Introduced concept
 - ▶ exhibited LSH for Hamming Distance.



- ▶ Locality Sensitive Hashing - introduced by Indyk & Motwani, in 1998.
 - ▶ Introduced concept
 - ▶ exhibited LSH for Hamming Distance.
- ▶ MinHash - by Broder, Charikar, Frieze, Mitzenmacher, 1998.
 - ▶ Introduced min-wise permutations.
 - ▶ LSH for Jaccard similarity.



- ▶ Locality Sensitive Hashing - introduced by Indyk & Motwani, in 1998.
 - ▶ Introduced concept
 - ▶ exhibited LSH for Hamming Distance.
- ▶ MinHash - by Broder, Charikar, Frieze, Mitzenmacher, 1998.
 - ▶ Introduced min-wise permutations.
 - ▶ LSH for Jaccard similarity.
- ▶ SimHash - by Charikar, 2002.
 - ▶ Demonstrated connections between randomized rounding and LSH
 - ▶ LSH for cosine similarity (angular distance).



Formal Definition

- ▶ A family of hash functions \mathcal{F} is a LSH if for *any* x, y the following holds:

$$\Pr_{h \in \mathcal{F}} [h(x) = h(y)] = \text{sim}(x, y)$$

where $\text{sim}(x, y)$ is a similarity measure.

- ▶ Read as: items that are “highly” similar land in the **same** hash-bucket with “high” probability, and items that are dissimilar land in the same hash bucket with “low” probability.



LSH for angular distance

- ▶ Given two vectors a, b , construct a hash function that maps these to the same bucket if the angle θ between them is *small*.
- ▶ Idea: use a random hyperplane (random projection).



LSH for angular distance

- ▶ Given two vectors a, b , construct a hash function that maps these to the same bucket if the angle θ between them is *small*.
- ▶ Idea: use a random hyperplane (random projection).



LSH for angular distance

- ▶ Given two vectors a, b , construct a hash function that maps these to the same bucket if the angle θ between them is *small*.
- ▶ Idea: use a random hyperplane (random projection).



LSH for angular distance

- ▶ This random hyperplane is a hash function $h : \mathbb{R}^n \rightarrow \{0, 1\}$.
- ▶ It maps a vector v to $\{0, 1\}$, depending on whether the vector v is to the top/bottom of the hyperplane (essentially, checking for the sign of the inner product).



Probability calculation

- ▶ Probability that vectors a, b map to the same output, i.e. $\Pr[h(a) = h(b)]?$



Probability calculation

- ▶ Probability that vectors a, b map to the same output, i.e. $\Pr[h(a) = h(b)]$?
- ▶ If θ be the angle between a, b , then this equals $\left[1 - \frac{\theta}{\pi}\right]$.



Precision? Recall?

- ▶ With a single hash function h , the precision may be quite low; the recall quite high.
- ▶ Just one hash may not be able to detect vectors that are indeed close by. So...



Precision? Recall?

- ▶ With a single hash function h , the precision may be quite low; the recall quite high.
- ▶ Just one hash may not be able to detect vectors that are indeed close by. So...
- ▶ Several (independent) copies of the hash function h ; call these h_1, h_2, \dots, h_k .



Precision? Recall?

- ▶ With k hash functions, we have a k -bit binary string, corresponding to a vector v .
- ▶ Call this the *hashcode* corresponding to the vector v (denoted as (v)).
- ▶ For two vectors a, b : if $(a) = (b)$ bit-by-bit, then surely a and b are close.



Precision? Recall?

- ▶ With k hash functions, we have a k -bit binary string, corresponding to a vector v .
- ▶ Call this the *hashcode* corresponding to the vector v (denoted as (v)).
- ▶ For two vectors a, b : if $(a) = (b)$ bit-by-bit, then surely a and b are close.
- ▶ This would help in improving precision. But recall may suffer.



Precision? Recall? Two extremes

- ▶ Given the hashcodes of the vectors, we can ask for a full bit-by-bit match to declare near duplicates. High precision, low recall.
- ▶ Given the hashcodes, we can ask for a *single* bit match to declare near duplicates. High recall, low precision.



Precision? Recall? Two extremes

- ▶ Given the hashcodes of the vectors, we can ask for a full bit-by-bit match to declare near duplicates. High precision, low recall.
- ▶ Given the hashcodes, we can ask for a *single* bit match to declare near duplicates. High recall, low precision.
- ▶ Mix the two: Banding. In the parlance of complexity theory, **gap amplification**.



Banding

- ▶ An example with $k = 15$ hashes, where `rows` = 3 and `bands` = 5.



Banding

- ▶ An example with $k = 15$ hashes, where `rows` = 3 and `bands` = 5.



Banding

- ▶ An example with $k = 15$ hashes, where `rows` = 3 and `bands` = 5.



Are a and b near-duplicates?

- ▶ We show the 's for a and b :



Are a and b near-duplicates?

- ▶ We show the 's for a and b :



Are a and b near-duplicates?

- ▶ We show the 's for a and b :



Are a and b near-duplicates?

- ▶ We declare a and b as near duplicates, if **there is a *band* in which they match bit-by-bit.**



Are a and b near-duplicates?

- ▶ We declare a and b as near duplicates, if **there is a *band* in which they match bit-by-bit.**
- ▶ To increase precision, increase the number of rows. To increase recall, increase number of bands.



Still a flourishing/active area of research

- ▶ Various considerations:
 - ▶ Engineering aspects: hash table construction time, query times, etc.
 - ▶ Engineering aspects: Maintain hashbuckets, update hashes. Bloom filters for hash buckets, etc.
 - ▶ How much randomness do I need?
 - ▶ How do I improve recall while maintaining precision.
 - ▶ Deep Hashing techniques?
 - ▶ Which similarity measures work best for different content: text, images, video.



Still a flourishing/active area of research

- ▶ Improve *recall* of LSH
 - ▶ covering LSH (only for Hamming space), (Pham-Pagh16)
 - ▶ other similarity measures wide open.
- ▶ Develop LSH's for other similarity measures.
 - ▶ Inner product (Neyshabur-Srebro15, Li-Shrivastava14)
 - ▶ Also gave rise to Assymmetric LSH.
- ▶ Improve training and query times based on data:
 - ▶ Data Dependent Hashing (Andoni-Razenshteyn15)
 - ▶ Learning to hash



LinkedIn's interest in finding near-duplicates

- ▶ Do we like memes on our LinkedIn page? Puzzles?
- ▶ Spam Filtering: spammers often use the same text repeatedly to *spam* members.



Thank You!



Thank You!

Questions?



Formal definition of LSH

- ▶ A hash is said to be a (S, cS, p_1, p_2) -LSH for a similarity function sim over the space \mathcal{X} if for any $x, y \in \mathcal{X}$:
 - ▶ if $\text{sim}(x, y) \geq S$ then $\Pr[h(x) = h(y)] \geq p_1$.
 - ▶ if $\text{sim}(x, y) \leq cS$ then $\Pr[h(x) = h(y)] \leq p_2$.

Here, $c \in (0, 1)$



Formal definition of LSH

- ▶ A hash is said to be a (S, cS, p_1, p_2) -LSH for a similarity function sim over the space \mathcal{X} if for any $x, y \in \mathcal{X}$:
 - ▶ if $\text{sim}(x, y) \geq S$ then $\Pr[h(x) = h(y)] \geq p_1$.
 - ▶ if $\text{sim}(x, y) \leq cS$ then $\Pr[h(x) = h(y)] \leq p_2$.

Here, $c \in (0, 1)$

- ▶ Read as: items that are “highly” similar land in the **same** hash-bucket with “high” probability, and items that are dissimilar land in the same hash bucket with “low” probability.
- ▶ Ideally we want: $p_1 = 1$, and $p_2 = 0$. These probabilities “mirror” the similarity function $\text{sim}(\cdot, \cdot)$.